

Staff AI Security Governance Engineer (AI GRC)

Role Overview

Coupang is seeking a Staff AI Security Governance Engineer to join our AI GRC (AI Governance, Risk Management, & Compliance) team in Seoul. This role is designed for a technical governance strategist who bridges the gap between complex AI regulatory requirements and technical security architecture.

In this role, you will define and architect Coupang's enterprise AI risk framework, translate regulatory mandates into policy-as-code specifications, and establish the technical guardrails for AI development. You will partner with AI/ML engineers, security architects, Legal, and Privacy teams to set standards, while collaborating with our dedicated Control Assurance team to ensure seamless execution and continuous monitoring.

What You Will Do

- **Framework Architecture & Policy-as-Code Strategy:** Design and maintain Coupang's AI Security & Risk Management Framework (mapped to NIST AI RMF, ISO 42001, MSIT guidelines, and global AI Acts). Architect policy-as-code rule sets and security baseline specifications for deployment across AI pipelines.
- **Assurance Partnering:** Define the technical control criteria and validation logic that the Control Assurance team will execute, monitor, and audit.
- **AI Risk Evaluation & Safeguards:** Identify, assess, and set mitigation standards for AI-specific threat vectors (e.g., prompt injection, training data leakage, model extraction, shadow AI, and supply chain vulnerabilities) across internal and third-party AI SaaS/LLM deployments.
- **Cross-Functional Engineering Alignment:** Partner closely with AI platform teams, infrastructure engineers, and security architects to embed governance and security requirements directly into AI architectures, CI/CD workflows, and runtime environments.
- **Enterprise AI Policy Ownership:** Own the enterprise AI policy lifecycle, establishing clear decision frameworks for AI model adoption, data privacy guardrails, and cross-border data transfer requirements.

Basic Qualifications

- Bachelor's degree in Computer Science, Information Security, Software Engineering, or a related technical/policy field.
- 10 + years of experience in Information Security, Technical GRC, Security Architecture, or AI/Cloud Governance.
- Deep understanding of AI/ML security risks, cloud security architecture (AWS), and privacy/data protection principles.

- Strong technical literacy to evaluate system architecture diagrams, interpret API/data flow models, and translate policy requirements into technical specifications for engineering and assurance teams.
- Proven track record of architecting governance frameworks (NIST AI RMF, ISO 42001, NIST CSF, PIPA) in tech-driven environments.
- Excellent written and verbal communication skills in English.

Preferred Qualifications

- Experience defining policy-as-code specifications, continuous control monitoring logic, or automated compliance guardrails in cloud environments.
- Hands-on familiarity with AI/ML pipelines, LLM integration architectures, data tokenization techniques, or cloud control planes.
- Experience partnering with operational audit/assurance teams to operationalize GRC metrics and controls.
- Relevant industry certifications (e.g., CISM, CRISC, CISSP, or specialized AI security/governance credentials).